

AWS EC2 GPU instance prices, published price list extract

Retrieved 29.08.2026, region us-east-1 (N. Virginia), Linux. Prepared as the §6.3 price documentation for the JouleFlex proposal (call 34/2026), subcontractor line: cloud GPU measurement runs.

Instance	GPU (count × model)	On-demand USD/hr	Spot USD/hr	On-demand USD per GPU-hr
g4dn.xlarge	1 × NVIDIA T4	0.526	0.220	0.526
g5.xlarge	1 × NVIDIA A10G	1.006	0.501	1.006
g6.xlarge	1 × NVIDIA L4	0.805	0.105	0.805
g6e.xlarge	1 × NVIDIA L40S	1.861	(not in public spot feed; per console)	1.861
p3.2xlarge	1 × NVIDIA V100	3.060	1.381	3.060
p4d.24xlarge	8 × NVIDIA A100 40GB	21.958	13.067	2.745
p5.48xlarge	8 × NVIDIA H100	55.040	57.762	6.880

Sources (AWS-published, publicly accessible).

On-demand: the pricing endpoint serving aws.amazon.com/ec2/pricing/on-demand/ — [bo.p.awsstatic.com/pricing/2.0/meteredUnitMaps/ec2/USD/current/ec2-ondemand-without-sec-sel/USEast \(N. Virginia\)/Linux/index.json](https://bo.p.awsstatic.com/pricing/2.0/meteredUnitMaps/ec2/USD/current/ec2-ondemand-without-sec-sel/USEast(N.Virginia)/Linux/index.json).

Spot: the AWS public spot price feed — spot-price.s3.amazonaws.com/spot.js (also rendered at aws.amazon.com/ec2/spot/pricing/).

Retrieved by automated fetch on 29.08.2026; values are AWS's published figures verbatim, per GPU-hour column computed by dividing the instance price by its GPU count. Spot prices vary continuously; the figures are the feed values at retrieval time. Companion files: full-page prints [aws-ec2-on-demand-2026-08-29.pdf](#) and [aws-ec2-spot-2026-08-29.pdf](#).

Budget relevance. The proposal plans 4,750 GPU-hours per year across card families, predominantly on single-GPU instances (T4, A10G, L4, L40S) with limited A100/H100 hours, using an on-demand/spot mix. At the prices above the planning figure of the budget specification is conservative.



Amazon EC2 ▾

[Pricing](#) > On-Demand Pricing

Amazon EC2 On-Demand Pricing



On-Demand Pricing

On-Demand Instances let you pay for compute capacity by the hour or second (minimum of 60 seconds) with no long-term commitments. This frees you from the costs and complexities of planning, purchasing, and maintaining hardware and transforms what are commonly large fixed costs into much smaller variable costs.



The pricing below includes the cost to run private and public AMIs on the specified operating system (“[Windows Usage](#)” prices apply to Windows Server 2003 R2, 2008, 2008 R2, 2012, 2012 R2, 2016, and 2019). Amazon also provides you with additional instances for [Amazon EC2 running Microsoft Windows with SQL Server](#), [Amazon EC2 running SUSE Linux Enterprise Server](#), and [Amazon EC2 running Red Hat Enterprise Linux](#).

Notice: Red Hat has made an update to their cloud pricing model for Red Hat Enterprise Linux (RHEL). On July 1, 2024 pricing for EC2 RHEL changed to a per-vCPU-hour based pricing model. Learn about the new prices in the [RHEL on AWS Pricing](#) page.

Chat Online

On-Demand Plans for Amazon EC2

Select a location type and region

Location Type

Region

AWS Region ▾

US East (Ohio) ▾

Select an operating system, instance type, and vCPU to view rates

Pricing is per instance-hour consumed for each instance, from the time an instance is launched until it is terminated or stopped. Each partial instance-hour consumed will be billed per-second for Linux, Windows, Windows with SQL Enterprise, Windows with SQL Standard, Windows with SQL Web, RHEL, and Ubuntu Pro Instances. SUSE Linux Enterprise Server (SLES) Instances are billed as a full hour.

T Amazon EC2 ▾

feature to modify the number of vCPUs and/or disable the Hyper-threading to save on vCPU-based licensing costs. The savings for using Windows Server AMIs is \$0.046 per vCPU-hour. The savings for using Windows with SQL Enterprise AMIs, Windows with SQL Standard AMIs, and Windows with SQL Web AMIs are \$0.421 per vCPU-hour, \$0.166 per vCPU-hour, and \$0.063 per vCPU-hour respectively. For EC2 instances with other operating systems, instance pricing will remain same as displayed above. For more details on valid vCPU counts and how to start using this feature, visit the [Optimize CPUs documentation page](#).

Looking for T1, M1, C1, CC2, M2, CR1, CG1, I2, HS1, M3, C3, or R3 instances? See the [Previous Generation Instances](#) page.

Data Transfer

The pricing below is based on data transferred "in" to and "out" of Amazon EC2.

If the Data Transfer per month is greater than 500 TB / month, please [contact us](#).



Rate tiers take into account your aggregate usage for Data Transfer Out to the Internet across Amazon EC2, Amazon S3, Amazon Glacier, Amazon RDS, Amazon Redshift, Amazon SageMaker, Amazon SES, Amazon SimpleDB, Amazon SQS, Amazon SNS, Amazon DynamoDB, AWS Storage Gateway, AWS CloudShell, and Amazon CloudWatch Logs.

AWS customers receive 100 GB of free data transfer out to the internet free each month, aggregated across all AWS Services and Regions (except China and GovCloud).

EU customers may request reduced data transfer rates for eligible use cases under the European Data Act. Please contact [AWS Customer Support](#) for more information.

AWS also offers eligible customers free data transfer out to the internet when they move all of their data off of AWS or all of their data off of a particular AWS service. Please check the [FAQ for more information](#), or [contact AWS Customer Support](#) for more information.

Data Transfer within the same AWS Region

Chat Online

Except as otherwise noted, our prices are exclusive of applicable taxes and duties, including VAT and applicable sales tax. For customers with a Japanese billing address, use of AWS is subject to Japanese Consumption Tax.

[Learn more about AWS tax policies.](#)

Amazon EC2 ▾

EBS-optimized instances enable EC2 instances to fully use the IOPS provisioned on an EBS volume. EBS-optimized instances deliver dedicated throughput between Amazon EC2 and Amazon EBS, with options between 500 and 4,000 Megabits per second (Mbps) depending on the instance type used. The dedicated throughput minimizes contention between Amazon EBS I/O and other traffic from your EC2 instance, providing the best performance for your EBS volumes. EBS-optimized instances are designed for use with both Standard and Provisioned IOPS Amazon EBS volumes. When attached to EBS-optimized instances, Provisioned IOPS volumes can achieve single digit millisecond latencies and are designed to deliver within 10% of the provisioned IOPS performance 99.9% of the time.

For Current Generation Instance types, EBS-optimization is enabled by default at no additional cost. For Previous Generation Instances types, EBS-optimization prices are on the Previous Generation Pricing Page.

The hourly price for EBS-optimized instances is in addition to the hourly usage fee for supported instance types.

Elastic IP Addresses

All In-use and Idle Elastic IP addresses are charged. To see public IPv4 address prices, visit the [VPC pricing page](#). 

Carrier IP Addresses

Except as otherwise noted, our prices are exclusive of applicable taxes and duties, including VAT and applicable sales tax. For customers with a Japanese billing address, use of AWS is subject to Japanese Consumption Tax.

[Learn more about AWS tax policies.](#)

Elastic Load Balancing

To see prices, visit the [Elastic Load Balancing](#) page.

On-Demand Capacity Reservations

On-Demand Capacity Reservations are priced exactly the same as their equivalent (On-Demand) instance usage. If a Capacity Reservation is fully utilized, you only pay for instance usage and nothing towards the Capacity Reservation. If a Capacity Reservation is partially utilized, you pay for the instance usage and for the unused portion of the Capacity Reservation. Learn more about [On-Demand Capacity Reservations](#).



Chat Online

T2/T3/T4g Unlimited Mode Pricing

For T4g instances in Unlimited mode, CPU Credits are charged at \$0.04 per vCPU-Hour for Linux, RHEL and SLES.

F Amazon EC2 ▼

- \$0.05 per vCPU-Hour for Linux, RHEL and SLES, and
- \$0.096 per vCPU-Hour for Windows and Windows with SQL Web

The CPU Credit pricing is the same for all instance sizes, for On-Demand, Spot, and Reserved Instances, and across all regions.

See [Unlimited Mode documentation](#) for details on when CPU Credits are charged.

Amazon CloudWatch

To see prices, visit the [Amazon Cloudwatch Pricing](#) page.

Amazon Elastic Block Store

To learn more, visit the [EBS Pricing Page](#).

Amazon EC2 Auto Scaling

Auto Scaling is enabled by Amazon CloudWatch and carries no additional fees. Each instance launched by Auto Scaling is automatically enabled for monitoring and the applicable [Amazon Cloudwatch](#) charges will be applied.

AWS GovCloud Region

[AWS GovCloud](#) is an AWS Region designed to allow U.S. government agencies and contractors to move more sensitive workloads into [the cloud](#) by addressing their specific regulatory and compliance requirements. For pricing and more information on the new AWS GovCloud Region, please visit the [AWS GovCloud Web Page](#).

** Your usage for the Free Tier is calculated each month across all regions except the AWS GovCloud region, and automatically applied to your bill – unused monthly usage will not roll over. Does not include Amazon EC2 running IBM, or the AWS GovCloud Region. See [offer terms](#) for more details and other restrictions.*

*** Rate tiers take into account your aggregate usage for Data Transfer Out to the Internet across Amazon EC2, Amazon S3, Amazon Glacier, Amazon RDS, Amazon Redshift, Amazon SageMaker, Amazon SES, Amazon SimpleDB, Amazon SQS, Amazon SNS, Amazon DynamoDB, AWS Storage Gateway, AWS CloudShell, and Amazon CloudWatch Logs.*

(Amazon EC2 is sold by Amazon Web Services, Inc.)



Chat Online

AMD SEV-SNP

Amazon EC2 supports AMD Secure Encrypted Virtualization-Secure Nested Paging (AMD SEV-SNP), a feature on AMD EPYC™ processors that provides enhanced security for your workloads. For more information, see [AMD SEV-SNP and Amazon EC2](#).

Amazon EC2 ▾

When you launch an Amazon EC2 instance with AMD SEV-SNP turned on, you are charged an additional hourly usage fee that is equivalent to 10 percent of the On-Demand hourly rate for the selected instance type.

This AMD SEV-SNP usage fee is a separate charge to your Amazon EC2 instance usage. Reserved Instances, Savings Plans, and operating system usage don't impact this fee

Additional pricing resources

[AWS Pricing Calculator](#)

Easily calculate your monthly costs with AWS

[Get pricing assistance](#)

Contact AWS specialists to get a personalized quote



Chat Online



Learn how to get started

There are several ways to get started with Amazon EC2 for free.

[Learn more about EC2](#)



Amazon EC2 ▾

Instantly get access to the AWS Free Tier.

[Sign up](#)



Start building in the console

Get started building with Amazon EC2 in the AWS Console.

[Sign in](#)

Reimagine Your Studio

Transform the way you work with virtual workstations on AWS



Optimize process manufacturing at scale

Converging OT and IT: Uncover predictive insights from time-series data



Chat Online

Proven Security to Accelerate Your Innovation

Find the resources you need to build, run, and scale your applications on the most secure cloud computing environment available today »



Amazon EC2 ▾

Learn



Resources



Developers



Help



Amazon is an equal opportunity employer and does not discriminate on the basis of protected veteran status, disability or other legally protected status. Veterans, military spouses, and people with disabilities are encouraged to apply.



Back to top ↑

Privacy Site terms Your Privacy Choices Cookie Preferences

© 2026, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Chat Online



Amazon EC2 ▾

[Spot Instances](#) > Pricing

Amazon EC2 Spot Instances Pricing



How does Amazon EC2 Spot Instances pricing work?

With Spot Instances, you pay the Spot price that's in effect for the time period your instances are running. Spot Instance prices are set by Amazon EC2 and adjust gradually based on long-term trends in supply and demand for Spot Instance capacity. To learn more about pricing, visit the [Spot Instance history page](#). The following table displays the Spot price for each region and instance type (updated every 5 minutes).



Spot Instances are available at a discount of up to 90% off compared to On-Demand pricing. To compare the current Spot prices against standard On-Demand rates, visit the [Spot Instance Advisor](#).

Note that you can now run Linux based RHEL and SUSE AMIs on Spot Instances. [Learn more](#).



Chat Online

Spot Instance Prices

Spot Instances



Note: T4g and T3 instances launch as unlimited by default. If you launch T4g or T3 Spot Instances as unlimited and plan to use them immediately and for a short duration, with no idle time for accruing CPU credits, you will incur charges for surplus credits. If the average CPU usage over a

24-hour period exceeds the baseline, you will also incur charges for surplus credits. We recommend that you launch your T4g or T3 Spot Instances in [standard mode](#) to avoid paying higher costs. For more information, see [Surplus Credits Can Incur Charges](#) and [Spot Instance limits](#).

Amazon EC2 ▾

Availability Zone, please refer to [Spot Instance pricing history](#) in the EC2 User Guide.

Notice: Red Hat has made an update to their cloud pricing model for Red Hat Enterprise Linux (RHEL). On July 1, 2024 pricing for EC2 RHEL changed to a per-vCPU-hour based pricing model. Learn about the new prices in the [RHEL on AWS Pricing](#) page.

Except as otherwise noted, our prices are exclusive of applicable taxes and duties, including VAT and applicable sales tax. For customers with a Japanese billing address, use of AWS is subject to Japanese Consumption Tax. [Learn more](#).



Additional pricing resources



Chat Online

Amazon FSx for Lustre



Accelerate compute workloads with a fast, parallel shared file system



Amazon EC2 ▾

Discover how AWS leads in development, deployment, and scalability »



Proven Security to Accelerate Your Innovation

Find the resources you need to build, run, and scale your applications on the most secure cloud computing environment available today »



Learn

Resources

Developers

Help


+ Chat Online



Amazon EC2 ▾

are encouraged to apply.



Back to top ↑

[Privacy](#) [Site terms](#) [Your Privacy Choices](#)  [Cookie Preferences](#)

© 2026, Amazon Web Services, Inc. or its affiliates. All rights reserved.



🕒 GPU DEDICATED SERVER HOSTING · ENTERPRISE GRADE

Rent GPU Dedicated Server — Enterprise Bare Metal GPU Server

Dedicated GPU server hosting for AI, HPC, rendering, and production workloads. Rent a bare metal GPU server with stable performance, dedicated resources, and reliable infrastructure.

[View GPU Server Pricing](#)

 Bare Metal GPU Server — No Sharing  Dedicated IPv4 + IPMI  Deploy in as fast as 10 minutes

 24/7 Expert Support  99.9% Uptime SLA

 GPU PLANS

Choose Your GPU Dedicated Server Plan

Rent dedicated server with GPU tailored to your workload. The longer your billing cycle, the lower the GPU dedicated server pricing. GTX, RTX, Quadro, and Data Center Nvidia GPU options available.

- All Plans
 - New Arrivals
 - Promotions
- GPU Use Scenario:
 - Live Streaming
 - HD Gaming
 - 3D Rendering
 - Video Editing
 - AI&Deep Learning
 - CAD/CGI/DCC
- GPU Memory:
 - 2 GB
 - 4 GB
 - 6 GB
 - 8 GB
 - 16 GB
 - 24 GB
 - 32 GB
 - 40 GB
 - 48 GB
 - 72 GB
 - 80 GB
 - 96 GB
 - 144 GB
 - 160 GB
 - 192 GB
- GPU Card Model:
 - P1000
 - GTX 1650
 - GTX 1660
 - RTX 2060
 - RTX 3060 Ti
 - RTX 4060
 - RTX 5060
 - RTX A4000
 - RTX A5000
 - RTX A6000
 - RTX 4090
 - RTX 5090
 - A100
 - H100

V100
P100
A40

Express Dedicated GPU Server - P1000

\$ 64.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: P1000
- CPU: 8-Core Xeon E5-2690
- Memory: 32GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 4 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - GTX 1650

\$ 99.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: GTX 1650
- CPU: 8-Core Xeon E5-2667v3
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 4 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - GTX 1660

\$ 139.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: GTX 1660
- CPU: 16-Core Dual E5-2660
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 6 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - RTX 4060

\$ 149.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 4060□
- CPU: 8-Core Xeon E5-2690
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 8 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - RTX 5060

\$ 159.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 5060□
- CPU: 24-Core Platinum 8160
- Memory: 64GB RAM
- Disk: 120GB SSD+960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 8 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Professional Dedicated GPU Server - P100

\$ 89.50/mo

55% OFF (Was \$199.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: P100□
- CPU: 16-Core Dual E5-2660
- Memory: 128GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Professional Dedicated GPU Server - RTX 2060

\$ 159.00/mo

20% OFF (Was \$199.00)

PrepaidOn-Demand

[Order Now](#)

- GPU Model: RTX 2060□
- CPU: 16-Core Dual E5-2660
- Memory: 128GB RAM

- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 6 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX 3060 Ti

\$ 179.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 3060 Ti☐
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 8 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - V100

\$ 229.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: V100☐
- CPU: 24-Core Dual E5-2690v3
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX 2060

\$ 143.40/mo

40% OFF (Was \$239.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 2060☐
- CPU: 40-Core Dual Gold 6148
- Memory: 128GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 6 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX A4000

\$ 209.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A4000□
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 16 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX A5000

\$ 164.03/mo

53% OFF (Was \$349.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A5000□
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 24 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX 4090

\$ 274.50/mo

50% OFF (Was \$549.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 4090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 24 GB GDDR6X

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A40

\$ 439.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: A40□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX 5090

\$ 479.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 5090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 32 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX A6000

\$ 409.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A6000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A100

\$ 359.55/mo

55% OFF (Was \$799.00)

PrepaidOn-Demand

[Order Now](#)

- GPU Model: A100□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 40 GB HBM2

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 2xRTX 5090

\$ 859.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 2 x RTX 5090□
- CPU: 44-core Dual E5-2699v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 32 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A100(80GB)

\$ 1559.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: A100(80GB)□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 80 GB HBM2e
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - H100

\$ 2099.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: H100□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 80 GB HBM2e
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 2xRTX 4090

\$ 431.52/mo

52% OFF (Was \$899.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 2 x RTX 4090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 24 GB GDDR6X
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xV100

\$ 469.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x V100□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xRTX A5000

\$ 539.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x RTX A5000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 24 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xRTX A6000

\$ 899.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x RTX A6000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 48 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 4xA100

\$ 1899.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 4 x A100 ☐
- CPU: 44-core Dual E5-2699v4
- Memory: 512GB RAM
- Disk: 240GB SSD+4TB NVMe+16TB SATA
- Bandwidth: 1000Mbps Unmetered ☐
- NVLink: 6xNVLink
- GPU Memory: 40 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 4xRTX A6000

\$ 1199.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 4 x RTX A6000 ☐
- CPU: 44-core Dual E5-2699v4
- Memory: 512GB RAM
- Disk: 240GB SSD+4TB NVMe+16TB SATA
- Bandwidth: 1000Mbps Unmetered ☐
- NVLink: 2xNVLink
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Support and Management Features for GPU Server



Supported Apps and Platforms for Dedicated GPU Server

Why Teams Choose Our GPU Dedicated Server Hosting

GPU dedicated server hosting that outperforms cloud instances — at a fraction of the cost. No virtualization overhead, full hardware control.



Dedicated Bare Metal GPU Server Performance

No virtualization overhead, noisy neighbors, or shared GPU contention. Every GPU, CPU, RAM, and storage resource is fully dedicated.



Multi-GPU & NVLink Ready

Scale from single-GPU deployments to multi-GPU configurations with NVLink support for large AI models and parallel workloads on your GPU dedicated server.



Large Local NVMe Storage Included

High-capacity NVMe and SSD storage options available without the expensive volume-based storage fees common on cloud GPU platforms.



Better Price-to-Performance Than Cloud

Predictable GPU dedicated server pricing with no GPU sharing, egress billing, or hidden infrastructure fees common on cloud platforms.



Enterprise Networking & Remote Access

Pure USA datacenter IPs, unmetered bandwidth, and IPMI remote



Flexible Configurations & Expert Support

Customize GPU, CPU, RAM, and storage to match your workload, backed by

management included with every GPU dedicated server.

experienced dedicated GPU server providers 24/7.

↗ VALUE COMPARISON

Dedicated GPU Server Cost Compared

Our cheap GPU dedicated server hosting delivers up to **3–5× better price-to-performance** than OVHCloud for equivalent GPU rentals. Consistently better than Hetzner after setup fees.

VRAM	PLATFORM	GPU	TIER	PRICE FROM	VALUE/\$ SCORE	VS GPU MART	NOTES
24GB	OVHCloud	L4	Mid-High	\$1,145/mo	0.00175	0.26× -74%	Storage extra; of out of st
24GB	Hetzner	RTX 4000	Low	\$214/mo	0.00467	0.70× -30%	\$500+ s fee
24GB	GPU Mart	A5000	High	\$269/mo	0.00669	1.0× Baseline ★	No hidd fees
48GB	OVHCloud	L40S	Very High	\$3,505/mo	0.00128	0.21× -79%	\$2,000 s fee
48GB	Contabo	L40S	Very High	\$802/mo	0.00561	0.88× -12%	Unstabl supply

VRAM	PLATFORM	GPU	TIER	PRICE FROM	VALUE/\$ SCORE	VS GPU MART	NOTES
48GB	GPU Mart	A6000	High	\$409/mo	0.00611	Best Value ★	Best price to-performance

Value/\$ = relative compute performance ÷ monthly price. Competitor pricing verified May 2026.

⚡ Rent GPU Dedicated Server Now

Compare GPU Dedicated Server Plans

USE CASES

Right GPU Dedicated Server for Your Workload

Match your GPU server rental to the right hardware — from RTX dedicated server options for rendering to A100 clusters for model training.

01 · AI INFERENCE & TRAINING

AI Model Serving, Training & Fine-Tuning

From LLM inference endpoints (vLLM, Ollama, TGI) to multi-week training runs on PyTorch and Hugging Face — a **GPU dedicated server** provides the sustained VRAM, full-node isolation, and NVLink multi-GPU support that shared cloud instances cannot match.

INFERENCE

RTX A4000 · A5000 · A40 · A6000

TRAINING

A100 40GB · A100 80GB · H100

[Explore AI & Deep Learning Servers →](#)

02 · STREAMING

Video Encoding & Live Streaming

Hardware NVENC acceleration with unmetered bandwidth and no encoding throttle. Suitable for OBS, FFmpeg, and 24/7 multi-stream pipelines on a **GTX or RTX dedicated server**.

RECOMMENDED GPUS

GTX 1650 · RTX 2060 · RTX 4060 · RTX 4090

[Explore Streaming Servers →](#)

03 · RENDERING

3D Rendering & Visual Production

Blender Cycles, Unreal Engine, and V-Ray scale directly with VRAM and CUDA cores. An **RTX dedicated server** with large VRAM handles high-poly scenes without CPU fallback — ideal for studios and render farms.

RECOMMENDED GPUS

RTX 3060 Ti · RTX A5000 · RTX A6000 · RTX 5090

[Explore Rendering Servers →](#)

GPU Dedicated Server Performance by Workload

Relative performance score (0–100) across 6 scenarios — helps you choose the right GPU server rental.

GPU	AI INFER	AI TRAIN	VIDEO ED.	RENDER	AUTO/SIM	GAMING
H100 80GB	100	100	75	80	100	25
RTX 5090	100	90	100	100	85	100
A100 40GB	90	95	65	70	95	20
RTX 4090	85	75	95	95	75	95

GPU	AI INFER	AI TRAIN	VIDEO ED.	RENDER	AUTO/SIM	GAMING
RTX A6000	75	70	85	95	90	75
A40	75	70	85	95	90	40
RTX 5060	65	45	75	70	45	75
RTX A5000	60	50	80	85	75	70
V100	50	65	40	45	65	15
RTX 4060	50	30	60	55	35	60
RTX A4000	45	35	70	70	60	55
RTX 3060 Ti	40	25	65	60	30	65
P100	30	45	30	25	45	10
RTX 2060	25	15	50	40	20	45
GTX 1660	12	6	40	18	8	30
T1000	10	5	25	15	10	15
GTX 1650	8	4	30	12	5	20

Excellent
Good
Moderate
Low
Basic

 QUICK DECISION GUIDE

Is a Dedicated GPU Server Hosting Plan Right for You?

GPU dedicated server hosting is designed for users who need consistent performance, isolated hardware, and predictable long-term GPU access.

✓ **Good fit if you need**

- Guaranteed access to dedicated GPU resources
- Stable monthly pricing without usage-based billing spikes
- 24/7 uptime on the same physical hardware
- Full root / administrator control over the environment
- Better isolation, privacy, and predictable performance

✗ **Probably not the best fit**

SITUATION	BETTER OPTION
Only need GPU occasionally	Hourly GPU
Very small budget or low monthly usage	GPU VPS
Need instant auto-scaling	Elastic cloud GPU
Prefer fully managed infrastructure	Managed AI platforms
Do not want server administration	Serverless GPU services
Looking for flexible GPU access? Explore our GPU VPS and hourly GPU servers .	

🖥️ **ARCHITECTURE**

GPU Dedicated Server **Architecture**

Built on company-owned physical hardware — not a cloud reseller. When you rent bare metal server hardware, you get full isolation, root access, and enterprise-grade I/O. Every dedicated server with GPU runs on its own node.


- **Exclusive Hardware Resources**

Exclusive CPU, GPU, memory, and NVMe on every node — zero noisy-neighbor contention, ever.
- **Flexible Customization**

When you rent dedicated server with GPU, any GPU, CPU, and storage combination is configured to your exact workload — no cloud markup.
- **Enterprise Network & Security**

IPv4/IPv6 dual-stack networking with optional Cisco ASA firewall, NVLink, HDMI dummy, and remote backup — included when you rent dedicated server with GPU.
- **Optimized I/O for GPU Workloads**

OS-level tuning with NVMe/SSD delivers low-latency, high-throughput I/O for AI and rendering pipelines on every GPU dedicated server.



☆ REVIEWS

What Our GPU Dedicated Server Customers Say

Real customers share their experience with our GPU server rental and dedicated GPU server hosting — verified on [Trustpilot](#)



"Easy to deal with. Well provisioned A100 server with plenty of RAM, CPU and disk."

James Freeman

Mar 31, 2026

☆ Trustpilot



"Extremely economical and great hosting service. Downtimes seldom happen and support is always there."

Haris Amjed

Mar 26, 2026

☆ Trustpilot

🕒 FAQ

GPU Dedicated Server FAQ

Everything about GPU server rental, cheap GPU dedicated server hosting plans, GTX & RTX dedicated server options, and how to rent GPU server resources for your workload.

GPU Dedicated Server FAQ Center

Which operating systems are available?

All major Linux distributions (Ubuntu, CentOS, Debian) and Windows Server are supported on our GPU server rental plans. Choose your preferred OS at checkout. For specific versions or custom configurations, contact our support team.

GeForce vs. Quadro vs. Tesla — which to rent?

A **GTX dedicated server** suits gaming, video transcoding, and entry-level HPC workloads at lowest cost. An **RTX dedicated server** adds Tensor Cores and ray-tracing for advanced AI inference and rendering. **Quadro (A4000/A5000/A6000)** offers professional-grade VRAM for 3D modeling, CAD, and AI. **Data Center (A100/H100)** maximizes FP32/FP64 and HBM memory for large-scale AI training and LLM serving.

Can I upgrade my GPU dedicated server?

Contact us with your requirements — our hardware team will confirm upgrade availability. When you rent dedicated server with GPU, you can upgrade disk, memory, bandwidth, and optional services. Switching GPU plans only charges the price difference. Note: adding more GPUs to an existing server is not supported.

How do I choose the right GPU for LLM inference?

For 7B–14B models, rent GPU server options like RTX A4000 or A5000 work well. Larger models, higher concurrency, or long-context inference benefit from an NVIDIA dedicated GPU server with A6000, A100, or H100. Choose based on model size, concurrent users, and latency requirements — not VRAM alone.

GPU VPS vs GPU Dedicated Server: which is better?

Our [GPU VPS](#) is ideal for testing, lightweight inference, and short-term GPU rental. A dedicated server with GPU provides fully isolated hardware for production AI inference, model training, rendering pipelines, and long-running workloads with consistent utilization.

Dedicated GPU vs Spot/Ephemeral GPU for production?

Spot GPUs reduce costs for fault-tolerant workloads but may be interrupted. Dedicated bare metal servers provide stable performance and uninterrupted access — better suited for production LLM serving, AI APIs, and latency-sensitive inference.

One large-VRAM GPU vs multiple smaller GPUs?

A single large-VRAM GPU is simpler for large models and long context windows. Multiple smaller GPUs can provide higher aggregate compute but require tensor parallelism. The best choice depends on whether your bottleneck is VRAM, throughput, or scalability. See our [multi-GPU server options](#).

Best GPU for AI image generation (Stable Diffusion, ComfyUI)?

RTX 4090, RTX 5090, A5000, and A6000 are popular for Stable Diffusion, Flux, and ComfyUI workflows on a GPU dedicated server. For high-throughput production pipelines or multi-user generation, an NVIDIA dedicated GPU server with A100 or H100 provides significantly higher concurrency and faster generation speed.

GPU dedicated server vs cloud GPU instances?

Hosting on dedicated hardware provides fully allocated resources with predictable performance and fixed monthly pricing. Cloud instances prioritize elasticity. When you rent bare metal server hardware for sustained workloads, it is typically more cost-efficient for AI inference, rendering, and long-running GPU server rental tasks.

Can I run multiple AI models on one GPU dedicated server?

Yes. Many customers run multiple inference models, embedding services, reranking models, or rendering workloads concurrently on a single GPU server, depending on available VRAM and concurrency requirements.

⚡ **READY TO GET STARTED?**

Rent a Nvidia GPU Dedicated Server

From a GTX dedicated server starting at \$64/mo to 4× A100 AI clusters — rent bare metal server hardware with transparent dedicated GPU server price options for every workload and budget.

⚡ **Rent GPU Dedicated Server**

All GPU Hosting Plans

