



Amazon EC2 ▾

[Pricing](#) > On-Demand Pricing

Amazon EC2 On-Demand Pricing



On-Demand Pricing

On-Demand Instances let you pay for compute capacity by the hour or second (minimum of 60 seconds) with no long-term commitments. This frees you from the costs and complexities of planning, purchasing, and maintaining hardware and transforms what are commonly large fixed costs into much smaller variable costs.



The pricing below includes the cost to run private and public AMIs on the specified operating system (“[Windows Usage](#)” prices apply to Windows Server 2003 R2, 2008, 2008 R2, 2012, 2012 R2, 2016, and 2019). Amazon also provides you with additional instances for [Amazon EC2 running Microsoft Windows with SQL Server](#), [Amazon EC2 running SUSE Linux Enterprise Server](#), and [Amazon EC2 running Red Hat Enterprise Linux](#).

Notice: Red Hat has made an update to their cloud pricing model for Red Hat Enterprise Linux (RHEL). On July 1, 2024 pricing for EC2 RHEL changed to a per-vCPU-hour based pricing model. Learn about the new prices in the [RHEL on AWS Pricing](#) page.

Chat Online

On-Demand Plans for Amazon EC2

Select a location type and region

Location Type

Region

AWS Region ▾

US East (Ohio) ▾

Select an operating system, instance type, and vCPU to view rates

Pricing is per instance-hour consumed for each instance, from the time an instance is launched until it is terminated or stopped. Each partial instance-hour consumed will be billed per-second for Linux, Windows, Windows with SQL Enterprise, Windows with SQL Standard, Windows with SQL Web, RHEL, and Ubuntu Pro Instances. SUSE Linux Enterprise Server (SLES) Instances are billed as a full hour.

T Amazon EC2 ▾

feature to modify the number of vCPUs and/or disable the Hyper-threading to save on vCPU-based licensing costs. The savings for using Windows Server AMIs is \$0.046 per vCPU-hour. The savings for using Windows with SQL Enterprise AMIs, Windows with SQL Standard AMIs, and Windows with SQL Web AMIs are \$0.421 per vCPU-hour, \$0.166 per vCPU-hour, and \$0.063 per vCPU-hour respectively. For EC2 instances with other operating systems, instance pricing will remain same as displayed above. For more details on valid vCPU counts and how to start using this feature, visit the [Optimize CPUs documentation page](#).

Looking for T1, M1, C1, CC2, M2, CR1, CG1, I2, HS1, M3, C3, or R3 instances? See the [Previous Generation Instances](#) page.

Data Transfer

The pricing below is based on data transferred "in" to and "out" of Amazon EC2.

If the Data Transfer per month is greater than 500 TB / month, please [contact us](#).



Rate tiers take into account your aggregate usage for Data Transfer Out to the Internet across Amazon EC2, Amazon S3, Amazon Glacier, Amazon RDS, Amazon Redshift, Amazon SageMaker, Amazon SES, Amazon SimpleDB, Amazon SQS, Amazon SNS, Amazon DynamoDB, AWS Storage Gateway, AWS CloudShell, and Amazon CloudWatch Logs.

AWS customers receive 100 GB of free data transfer out to the internet free each month, aggregated across all AWS Services and Regions (except China and GovCloud).

EU customers may request reduced data transfer rates for eligible use cases under the European Data Act. Please contact [AWS Customer Support](#) for more information.

AWS also offers eligible customers free data transfer out to the internet when they move all of their data off of AWS or all of their data off of a particular AWS service. Please check the [FAQ for more information](#), or [contact AWS Customer Support](#) for more information.

Data Transfer within the same AWS Region

Chat Online

Except as otherwise noted, our prices are exclusive of applicable taxes and duties, including VAT and applicable sales tax. For customers with a Japanese billing address, use of AWS is subject to Japanese Consumption Tax.

[Learn more about AWS tax policies.](#)

Amazon EC2 ▾

EBS-optimized instances enable EC2 instances to fully use the IOPS provisioned on an EBS volume. EBS-optimized instances deliver dedicated throughput between Amazon EC2 and Amazon EBS, with options between 500 and 4,000 Megabits per second (Mbps) depending on the instance type used. The dedicated throughput minimizes contention between Amazon EBS I/O and other traffic from your EC2 instance, providing the best performance for your EBS volumes. EBS-optimized instances are designed for use with both Standard and Provisioned IOPS Amazon EBS volumes. When attached to EBS-optimized instances, Provisioned IOPS volumes can achieve single digit millisecond latencies and are designed to deliver within 10% of the provisioned IOPS performance 99.9% of the time.

For Current Generation Instance types, EBS-optimization is enabled by default at no additional cost. For Previous Generation Instances types, EBS-optimization prices are on the Previous Generation Pricing Page.

The hourly price for EBS-optimized instances is in addition to the hourly usage fee for supported instance types.

Elastic IP Addresses

All In-use and Idle Elastic IP addresses are charged. To see public IPv4 address prices, visit the [VPC pricing page](#). 

Carrier IP Addresses

Except as otherwise noted, our prices are exclusive of applicable taxes and duties, including VAT and applicable sales tax. For customers with a Japanese billing address, use of AWS is subject to Japanese Consumption Tax.

[Learn more about AWS tax policies.](#)

Elastic Load Balancing

To see prices, visit the [Elastic Load Balancing](#) page.

On-Demand Capacity Reservations

On-Demand Capacity Reservations are priced exactly the same as their equivalent (On-Demand) instance usage. If a Capacity Reservation is fully utilized, you only pay for instance usage and nothing towards the Capacity Reservation. If a Capacity Reservation is partially utilized, you pay for the instance usage and for the unused portion of the Capacity Reservation. Learn more about [On-Demand Capacity Reservations](#).



Chat Online

T2/T3/T4g Unlimited Mode Pricing

For T4g instances in Unlimited mode, CPU Credits are charged at \$0.04 per vCPU-Hour for Linux, RHEL and SLES.

F Amazon EC2 ▾

- \$0.05 per vCPU-Hour for Linux, RHEL and SLES, and
- \$0.096 per vCPU-Hour for Windows and Windows with SQL Web

The CPU Credit pricing is the same for all instance sizes, for On-Demand, Spot, and Reserved Instances, and across all regions.

See [Unlimited Mode documentation](#) for details on when CPU Credits are charged.

Amazon CloudWatch

To see prices, visit the [Amazon Cloudwatch Pricing](#) page.

Amazon Elastic Block Store

To learn more, visit the [EBS Pricing Page](#).

Amazon EC2 Auto Scaling

Auto Scaling is enabled by Amazon CloudWatch and carries no additional fees. Each instance launched by Auto Scaling is automatically enabled for monitoring and the applicable [Amazon Cloudwatch](#) charges will be applied.

AWS GovCloud Region

[AWS GovCloud](#) is an AWS Region designed to allow U.S. government agencies and contractors to move more sensitive workloads into [the cloud](#) by addressing their specific regulatory and compliance requirements. For pricing and more information on the new AWS GovCloud Region, please visit the [AWS GovCloud Web Page](#).

** Your usage for the Free Tier is calculated each month across all regions except the AWS GovCloud region, and automatically applied to your bill – unused monthly usage will not roll over. Does not include Amazon EC2 running IBM, or the AWS GovCloud Region. See [offer terms](#) for more details and other restrictions.*

*** Rate tiers take into account your aggregate usage for Data Transfer Out to the Internet across Amazon EC2, Amazon S3, Amazon Glacier, Amazon RDS, Amazon Redshift, Amazon SageMaker, Amazon SES, Amazon SimpleDB, Amazon SQS, Amazon SNS, Amazon DynamoDB, AWS Storage Gateway, AWS CloudShell, and Amazon CloudWatch Logs.*

(Amazon EC2 is sold by Amazon Web Services, Inc.)



Chat Online

AMD SEV-SNP

Amazon EC2 supports AMD Secure Encrypted Virtualization-Secure Nested Paging (AMD SEV-SNP), a feature on AMD EPYC™ processors that provides enhanced security for your workloads. For more information, see [AMD SEV-SNP and Amazon EC2](#).

Amazon EC2 ▼

When you launch an Amazon EC2 instance with AMD SEV-SNP turned on, you are charged an additional hourly usage fee that is equivalent to 10 percent of the On-Demand hourly rate for the selected instance type.

This AMD SEV-SNP usage fee is a separate charge to your Amazon EC2 instance usage. Reserved Instances, Savings Plans, and operating system usage don't impact this fee

Additional pricing resources

[AWS Pricing Calculator](#)

Easily calculate your monthly costs with AWS

[Get pricing assistance](#)

Contact AWS specialists to get a personalized quote



Chat Online



Learn how to get started

There are several ways to get started with Amazon EC2 for free.

[Learn more about EC2](#)



Amazon EC2 ▾

Instantly get access to the AWS Free Tier.

[Sign up](#)



Start building in the console

Get started building with Amazon EC2 in the AWS Console.

[Sign in](#)

Reimagine Your Studio

Transform the way you work with virtual workstations on AWS



Optimize process manufacturing at scale

Converging OT and IT: Uncover predictive insights from time-series data



Chat Online

Proven Security to Accelerate Your Innovation

Find the resources you need to build, run, and scale your applications on the most secure cloud computing environment available today »



Amazon EC2 ▾

Learn



Resources



Developers



Help



Amazon is an equal opportunity employer and does not discriminate on the basis of protected veteran status, disability or other legally protected status. Veterans, military spouses, and people with disabilities are encouraged to apply.



Back to top ↑

Privacy Site terms Your Privacy Choices Cookie Preferences

© 2026, Amazon Web Services, Inc. or its affiliates. All rights reserved.



Chat Online