

AWS EC2 GPU instance prices, published price list extract

Retrieved 29.08.2026, region us-east-1 (N. Virginia), Linux. Prepared as the §6.3 price documentation for the JouleFlex proposal (call 34/2026), subcontractor line: cloud GPU measurement runs.

Instance	GPU (count × model)	On-demand USD/hr	Spot USD/hr	On-demand USD per GPU-hr
g4dn.xlarge	1 × NVIDIA T4	0.526	0.220	0.526
g5.xlarge	1 × NVIDIA A10G	1.006	0.501	1.006
g6.xlarge	1 × NVIDIA L4	0.805	0.105	0.805
g6e.xlarge	1 × NVIDIA L40S	1.861	(not in public spot feed; per console)	1.861
p3.2xlarge	1 × NVIDIA V100	3.060	1.381	3.060
p4d.24xlarge	8 × NVIDIA A100 40GB	21.958	13.067	2.745
p5.48xlarge	8 × NVIDIA H100	55.040	57.762	6.880

Sources (AWS-published, publicly accessible).

On-demand: the pricing endpoint serving aws.amazon.com/ec2/pricing/on-demand/ — [bo.p.awsstatic.com/pricing/2.0/meteredUnitMaps/ec2/USD/current/ec2-ondemand-without-sec-sel/USEast \(N. Virginia\)/Linux/index.json](https://bo.p.awsstatic.com/pricing/2.0/meteredUnitMaps/ec2/USD/current/ec2-ondemand-without-sec-sel/USEast(N.Virginia)/Linux/index.json).

Spot: the AWS public spot price feed — spot-price.s3.amazonaws.com/spot.js (also rendered at aws.amazon.com/ec2/spot/pricing/).

Retrieved by automated fetch on 29.08.2026; values are AWS's published figures verbatim, per GPU-hour column computed by dividing the instance price by its GPU count. Spot prices vary continuously; the figures are the feed values at retrieval time. Companion files: full-page prints [aws-ec2-on-demand-2026-08-29.pdf](#) and [aws-ec2-spot-2026-08-29.pdf](#).

Budget relevance. The proposal plans 4,750 GPU-hours per year across card families, predominantly on single-GPU instances (T4, A10G, L4, L40S) with limited A100/H100 hours, using an on-demand/spot mix. At the prices above the planning figure of the budget specification is conservative.