

🕒 GPU DEDICATED SERVER HOSTING · ENTERPRISE GRADE

Rent GPU Dedicated Server — Enterprise Bare Metal GPU Server

Dedicated GPU server hosting for AI, HPC, rendering, and production workloads. Rent a bare metal GPU server with stable performance, dedicated resources, and reliable infrastructure.

[View GPU Server Pricing](#)

-  Bare Metal GPU Server — No Sharing
-  Dedicated IPv4 + IPMI
-  Deploy in as fast as 10 minutes
-  24/7 Expert Support
-  99.9% Uptime SLA

 GPU PLANS

Choose Your GPU Dedicated Server Plan

Rent dedicated server with GPU tailored to your workload. The longer your billing cycle, the lower the GPU dedicated server pricing. GTX, RTX, Quadro, and Data Center Nvidia GPU options available.

- All Plans
 - New Arrivals
 - Promotions
- GPU Use Scenario:
 - Live Streaming
 - HD Gaming
 - 3D Rendering
 - Video Editing
 - AI&Deep Learning
 - CAD/CGI/DCC
- GPU Memory:
 - 2 GB
 - 4 GB
 - 6 GB
 - 8 GB
 - 16 GB
 - 24 GB
 - 32 GB
 - 40 GB
 - 48 GB
 - 72 GB
 - 80 GB
 - 96 GB
 - 144 GB
 - 160 GB
 - 192 GB
- GPU Card Model:
 - P1000
 - GTX 1650
 - GTX 1660
 - RTX 2060
 - RTX 3060 Ti
 - RTX 4060
 - RTX 5060
 - RTX A4000
 - RTX A5000
 - RTX A6000
 - RTX 4090
 - RTX 5090
 - A100
 - H100

V100
P100
A40

Express Dedicated GPU Server - P1000

\$ 64.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: P1000
- CPU: 8-Core Xeon E5-2690
- Memory: 32GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 4 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - GTX 1650

\$ 99.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: GTX 1650
- CPU: 8-Core Xeon E5-2667v3
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 4 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - GTX 1660

\$ 139.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: GTX 1660
- CPU: 16-Core Dual E5-2660
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 6 GB GDDR5
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - RTX 4060

\$ 149.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 4060□
- CPU: 8-Core Xeon E5-2690
- Memory: 64GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 8 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Basic Dedicated GPU Server - RTX 5060

\$ 159.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 5060□
- CPU: 24-Core Platinum 8160
- Memory: 64GB RAM
- Disk: 120GB SSD+960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 8 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Professional Dedicated GPU Server - P100

\$ 89.50/mo

55% OFF (Was \$199.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: P100□
- CPU: 16-Core Dual E5-2660
- Memory: 128GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Professional Dedicated GPU Server - RTX 2060

\$ 159.00/mo

20% OFF (Was \$199.00)

PrepaidOn-Demand

[Order Now](#)

- GPU Model: RTX 2060□
- CPU: 16-Core Dual E5-2660
- Memory: 128GB RAM

- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 6 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX 3060 Ti

\$ 179.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 3060 Ti☐
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 8 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - V100

\$ 229.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: V100☐
- CPU: 24-Core Dual E5-2690v3
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX 2060

\$ 143.40/mo

40% OFF (Was \$239.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 2060☐
- CPU: 40-Core Dual Gold 6148
- Memory: 128GB RAM
- Disk: 120GB SSD + 960GB SSD
- Bandwidth: 100Mbps Unmetered☐
- GPU Memory: 6 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX A4000

\$ 209.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A4000□
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 16 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Advanced Dedicated GPU Server - RTX A5000

\$ 164.03/mo

53% OFF (Was \$349.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A5000□
- CPU: 24-Core Dual E5-2697v2
- Memory: 128GB RAM
- Disk: 240GB SSD+2TB SSD
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 24 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX 4090

\$ 274.50/mo

50% OFF (Was \$549.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 4090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 24 GB GDDR6X

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A40

\$ 439.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: A40□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX 5090

\$ 479.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX 5090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 32 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - RTX A6000

\$ 409.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: RTX A6000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A100

\$ 359.55/mo

55% OFF (Was \$799.00)

PrepaidOn-Demand

[Order Now](#)

- GPU Model: A100□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered□
- GPU Memory: 40 GB HBM2

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 2xRTX 5090

\$ 859.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 2 x RTX 5090
- CPU: 44-core Dual E5-2699v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered
- GPU Memory: 32 GB GDDR7
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - A100(80GB)

\$ 1559.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: A100(80GB)
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 80 GB HBM2e
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Dedicated GPU Server - H100

\$ 2099.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: H100
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 100Mbps Unmetered
- GPU Memory: 80 GB HBM2e
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 2xRTX 4090

\$ 431.52/mo

52% OFF (Was \$899.00)

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 2 x RTX 4090□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 24 GB GDDR6X
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xV100

\$ 469.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x V100□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 16 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xRTX A5000

\$ 539.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x RTX A5000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 24 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 3xRTX A6000

\$ 899.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 3 x RTX A6000□
- CPU: 36-Core Dual E5-2697v4
- Memory: 256GB RAM
- Disk: 240GB SSD+2TB NVMe+8TB SATA
- Bandwidth: 1000Mbps Unmetered□
- GPU Memory: 48 GB GDDR6

- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 4xA100

\$ 1899.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 4 x A100 ☐
- CPU: 44-core Dual E5-2699v4
- Memory: 512GB RAM
- Disk: 240GB SSD+4TB NVMe+16TB SATA
- Bandwidth: 1000Mbps Unmetered ☐
- NVLink: 6xNVLink
- GPU Memory: 40 GB HBM2
- IP: 1 Dedicated IPv4
- Location: USA

Enterprise Multi-GPU Dedicated Server - 4xRTX A6000

\$ 1199.00/mo

1mo3mo12mo24mo

[Order Now](#)

- GPU Model: 4 x RTX A6000 ☐
- CPU: 44-core Dual E5-2699v4
- Memory: 512GB RAM
- Disk: 240GB SSD+4TB NVMe+16TB SATA
- Bandwidth: 1000Mbps Unmetered ☐
- NVLink: 2xNVLink
- GPU Memory: 48 GB GDDR6
- IP: 1 Dedicated IPv4
- Location: USA

Support and Management Features for GPU Server



widgetsSupported Apps and Platforms for Dedicated GPU Server

Why Teams Choose Our GPU Dedicated Server Hosting

GPU dedicated server hosting that outperforms cloud instances — at a fraction of the cost. No virtualization overhead, full hardware control.



Dedicated Bare Metal GPU Server Performance

No virtualization overhead, noisy neighbors, or shared GPU contention. Every GPU, CPU, RAM, and storage resource is fully dedicated.



Multi-GPU & NVLink Ready

Scale from single-GPU deployments to multi-GPU configurations with NVLink support for large AI models and parallel workloads on your GPU dedicated server.



Large Local NVMe Storage Included

High-capacity NVMe and SSD storage options available without the expensive volume-based storage fees common on cloud GPU platforms.



Better Price-to-Performance Than Cloud

Predictable GPU dedicated server pricing with no GPU sharing, egress billing, or hidden infrastructure fees common on cloud platforms.



Enterprise Networking & Remote Access

Pure USA datacenter IPs, unmetered bandwidth, and IPMI remote



Flexible Configurations & Expert Support

Customize GPU, CPU, RAM, and storage to match your workload, backed by

management included with every GPU dedicated server.

experienced dedicated GPU server providers 24/7.

↗ VALUE COMPARISON

Dedicated GPU Server Cost Compared

Our cheap GPU dedicated server hosting delivers up to **3–5× better price-to-performance** than OVHCloud for equivalent GPU rentals. Consistently better than Hetzner after setup fees.

VRAM	PLATFORM	GPU	TIER	PRICE FROM	VALUE/\$ SCORE	VS GPU MART	NOTES
24GB	OVHCloud	L4	Mid-High	\$1,145/mo	0.00175	0.26× -74%	Storage extra; of out of st
24GB	Hetzner	RTX 4000	Low	\$214/mo	0.00467	0.70× -30%	\$500+ s fee
24GB	GPU Mart	A5000	High	\$269/mo	0.00669	1.0× Baseline ★	No hidd fees
48GB	OVHCloud	L40S	Very High	\$3,505/mo	0.00128	0.21× -79%	\$2,000 s fee
48GB	Contabo	L40S	Very High	\$802/mo	0.00561	0.88× -12%	Unstabl supply

VRAM	PLATFORM	GPU	TIER	PRICE FROM	VALUE/\$ SCORE	VS GPU MART	NOTES
48GB	GPU Mart	A6000	High	\$409/mo	0.00611	Best Value ★	Best price to-performance

Value/\$ = relative compute performance ÷ monthly price. Competitor pricing verified May 2026.

⚡ Rent GPU Dedicated Server Now

Compare GPU Dedicated Server Plans

USE CASES

Right GPU Dedicated Server for Your Workload

Match your GPU server rental to the right hardware — from RTX dedicated server options for rendering to A100 clusters for model training.

01 · AI INFERENCE & TRAINING

AI Model Serving, Training & Fine-Tuning

From LLM inference endpoints (vLLM, Ollama, TGI) to multi-week training runs on PyTorch and Hugging Face — a **GPU dedicated server** provides the sustained VRAM, full-node isolation, and NVLink multi-GPU support that shared cloud instances cannot match.

INFERENCE

RTX A4000 · A5000 · A40 · A6000

TRAINING

A100 40GB · A100 80GB · H100

[Explore AI & Deep Learning Servers →](#)

02 · STREAMING

Video Encoding & Live Streaming

Hardware NVENC acceleration with unmetered bandwidth and no encoding throttle. Suitable for OBS, FFmpeg, and 24/7 multi-stream pipelines on a **GTX or RTX dedicated server**.

RECOMMENDED GPUS

GTX 1650 · RTX 2060 · RTX 4060 · RTX 4090

[Explore Streaming Servers →](#)

03 · RENDERING

3D Rendering & Visual Production

Blender Cycles, Unreal Engine, and V-Ray scale directly with VRAM and CUDA cores. An **RTX dedicated server** with large VRAM handles high-poly scenes without CPU fallback — ideal for studios and render farms.

RECOMMENDED GPUS

RTX 3060 Ti · RTX A5000 · RTX A6000 · RTX 5090

[Explore Rendering Servers →](#)

GPU Dedicated Server Performance by Workload

Relative performance score (0–100) across 6 scenarios — helps you choose the right GPU server rental.

GPU	AI INFER	AI TRAIN	VIDEO ED.	RENDER	AUTO/SIM	GAMING
H100 80GB	100	100	75	80	100	25
RTX 5090	100	90	100	100	85	100
A100 40GB	90	95	65	70	95	20
RTX 4090	85	75	95	95	75	95

GPU	AI INFER	AI TRAIN	VIDEO ED.	RENDER	AUTO/SIM	GAMING
RTX A6000	75	70	85	95	90	75
A40	75	70	85	95	90	40
RTX 5060	65	45	75	70	45	75
RTX A5000	60	50	80	85	75	70
V100	50	65	40	45	65	15
RTX 4060	50	30	60	55	35	60
RTX A4000	45	35	70	70	60	55
RTX 3060 Ti	40	25	65	60	30	65
P100	30	45	30	25	45	10
RTX 2060	25	15	50	40	20	45
GTX 1660	12	6	40	18	8	30
T1000	10	5	25	15	10	15
GTX 1650	8	4	30	12	5	20

Excellent

Good

Moderate

Low

Basic

 QUICK DECISION GUIDE

Is a Dedicated GPU Server Hosting Plan Right for You?

GPU dedicated server hosting is designed for users who need consistent performance, isolated hardware, and predictable long-term GPU access.

✓ **Good fit if you need**

- Guaranteed access to dedicated GPU resources
- Stable monthly pricing without usage-based billing spikes
- 24/7 uptime on the same physical hardware
- Full root / administrator control over the environment
- Better isolation, privacy, and predictable performance

✗ **Probably not the best fit**

SITUATION	BETTER OPTION
Only need GPU occasionally	Hourly GPU
Very small budget or low monthly usage	GPU VPS
Need instant auto-scaling	Elastic cloud GPU
Prefer fully managed infrastructure	Managed AI platforms
Do not want server administration	Serverless GPU services
Looking for flexible GPU access? Explore our GPU VPS and hourly GPU servers .	

 ARCHITECTURE

GPU Dedicated Server **Architecture**

Built on company-owned physical hardware — not a cloud reseller. When you rent bare metal server hardware, you get full isolation, root access, and enterprise-grade I/O. Every dedicated server with GPU runs on its own node.


- **Exclusive Hardware Resources**

Exclusive CPU, GPU, memory, and NVMe on every node — zero noisy-neighbor contention, ever.
- **Flexible Customization**

When you rent dedicated server with GPU, any GPU, CPU, and storage combination is configured to your exact workload — no cloud markup.
- **Enterprise Network & Security**

IPv4/IPv6 dual-stack networking with optional Cisco ASA firewall, NVLink, HDMI dummy, and remote backup — included when you rent dedicated server with GPU.
- **Optimized I/O for GPU Workloads**

OS-level tuning with NVMe/SSD delivers low-latency, high-throughput I/O for AI and rendering pipelines on every GPU dedicated server.



☆ REVIEWS

What Our GPU Dedicated Server Customers Say

Real customers share their experience with our GPU server rental and dedicated GPU server hosting — verified on [Trustpilot](#)



"Easy to deal with. Well provisioned A100 server with plenty of RAM, CPU and disk."

James Freeman

Mar 31, 2026

☆ Trustpilot



"Extremely economical and great hosting service. Downtimes seldom happen and support is always there."

Haris Amjed

Mar 26, 2026

☆ Trustpilot

🕒 FAQ

GPU Dedicated Server FAQ

Everything about GPU server rental, cheap GPU dedicated server hosting plans, GTX & RTX dedicated server options, and how to rent GPU server resources for your workload.

GPU Dedicated Server FAQ Center

Which operating systems are available?

All major Linux distributions (Ubuntu, CentOS, Debian) and Windows Server are supported on our GPU server rental plans. Choose your preferred OS at checkout. For specific versions or custom configurations, contact our support team.

GeForce vs. Quadro vs. Tesla — which to rent?

A **GTX dedicated server** suits gaming, video transcoding, and entry-level HPC workloads at lowest cost. An **RTX dedicated server** adds Tensor Cores and ray-tracing for advanced AI inference and rendering. **Quadro (A4000/A5000/A6000)** offers professional-grade VRAM for 3D modeling, CAD, and AI. **Data Center (A100/H100)** maximizes FP32/FP64 and HBM memory for large-scale AI training and LLM serving.

Can I upgrade my GPU dedicated server?

Contact us with your requirements — our hardware team will confirm upgrade availability. When you rent dedicated server with GPU, you can upgrade disk, memory, bandwidth, and optional services. Switching GPU plans only charges the price difference. Note: adding more GPUs to an existing server is not supported.

How do I choose the right GPU for LLM inference?

For 7B–14B models, rent GPU server options like RTX A4000 or A5000 work well. Larger models, higher concurrency, or long-context inference benefit from an NVIDIA dedicated GPU server with A6000, A100, or H100. Choose based on model size, concurrent users, and latency requirements — not VRAM alone.

GPU VPS vs GPU Dedicated Server: which is better?

Our [GPU VPS](#) is ideal for testing, lightweight inference, and short-term GPU rental. A dedicated server with GPU provides fully isolated hardware for production AI inference, model training, rendering pipelines, and long-running workloads with consistent utilization.

Dedicated GPU vs Spot/Ephemeral GPU for production?

Spot GPUs reduce costs for fault-tolerant workloads but may be interrupted. Dedicated bare metal servers provide stable performance and uninterrupted access — better suited for production LLM serving, AI APIs, and latency-sensitive inference.

One large-VRAM GPU vs multiple smaller GPUs?

A single large-VRAM GPU is simpler for large models and long context windows. Multiple smaller GPUs can provide higher aggregate compute but require tensor parallelism. The best choice depends on whether your bottleneck is VRAM, throughput, or scalability. See our [multi-GPU server options](#).

Best GPU for AI image generation (Stable Diffusion, ComfyUI)?

RTX 4090, RTX 5090, A5000, and A6000 are popular for Stable Diffusion, Flux, and ComfyUI workflows on a GPU dedicated server. For high-throughput production pipelines or multi-user generation, an NVIDIA dedicated GPU server with A100 or H100 provides significantly higher concurrency and faster generation speed.

GPU dedicated server vs cloud GPU instances?

Hosting on dedicated hardware provides fully allocated resources with predictable performance and fixed monthly pricing. Cloud instances prioritize elasticity. When you rent bare metal server hardware for sustained workloads, it is typically more cost-efficient for AI inference, rendering, and long-running GPU server rental tasks.

Can I run multiple AI models on one GPU dedicated server?

Yes. Many customers run multiple inference models, embedding services, reranking models, or rendering workloads concurrently on a single GPU server, depending on available VRAM and concurrency requirements.

⚡ **READY TO GET STARTED?**

Rent a Nvidia GPU Dedicated Server

From a GTX dedicated server starting at \$64/mo to 4× A100 AI clusters — rent bare metal server hardware with transparent dedicated GPU server price options for every workload and budget.

⚡ **Rent GPU Dedicated Server**

All GPU Hosting Plans

